

Predicting mortality caused by heart failure

Anton Ivanov

2021-12-09

1 Introduction

In this project we will train a model to predict mortality caused by heart failure (HF) based on patients features. Acute heart failure is a leading cause of hospitalization and death, and it is an increasing burden on health care systems. The correct risk stratification of patients could improve clinical outcome and resources allocation, avoiding the overtreatment of low-risk subjects or the early, inappropriate discharge of high-risk patients.¹

Heart failure occurs when the heart becomes too weak or stiff to pump enough blood to meet the body's needs. Symptoms typically include breathlessness, extreme fatigue, reduced capacity to exercise, etc. The number of people living with HF is high and growing. More than 15 million people (~ 2%) are estimated to be living with HF in Europe. People living with HF are at high risk of hospitalisation, despite improvements in treatment options and care in the past two decades, mortality from HF remains high.²

The used dataset contains medical records of 299 patients suffering from heart failure, collected during their follow-up period. Each patient profile includes 11 clinical features including blood measurements, health indicators and other relevant information. The dataset can be downloaded from the [UCI Machine Learning Repository] (<https://archive-beta.ics.uci.edu/ml/datasets/heart+failure+clinical+records>).

The goal of this project is to develop a machine learning model that will provide the best prediction of the patients death based on the available information. *caret* package will be used for training and assessing the models. Model performance will be evaluated using a metric most appropriate for the current data.

2 Analysis and methods

2.1 Data preparation

The dataset is available as a CSV file with columns separated by commas. It is saved as a data frame which has the following structure:

```
data <- read.csv("heart_failure_clinical_records_dataset.csv")
str(data)
```

```
## 'data.frame':    299 obs. of  13 variables:
## $ age           : num  75 55 65 50 65 90 75 60 65 80 ...
## $ anaemia       : int   0 0 0 1 1 1 1 0 1 ...
## $ creatinine_phosphokinase: int 582 7861 146 111 160 47 246 315 157 123 ...
## $ diabetes      : int   0 0 0 0 1 0 0 1 0 0 ...
```

¹World J Cardiol. 2015 Dec 26; 7(12): 902–911.

²Heart Failure Policy Network. 2020. Heart failure policy and practice in Europe. London: HFPN

```
## $ ejection_fraction      : int  20 38 20 20 20 40 15 60 65 35 ...
## $ high_blood_pressure    : int   1 0 0 0 0 1 0 0 0 1 ...
## $ platelets              : num 265000 263358 162000 210000 327000 ...
## $ serum_creatinine       : num  1.9 1.1 1.3 1.9 2.7 2.1 1.2 1.1 1.5 9.4 ...
## $ serum_sodium           : int  130 136 129 137 116 132 137 131 138 133 ...
## $ sex                    : int   1 1 1 1 0 1 1 1 0 1 ...
## $ smoking                : int   0 0 1 0 0 1 0 1 0 1 ...
## $ time                   : int   4 6 7 7 8 8 10 10 10 10 ...
## $ DEATH_EVENT            : int   1 1 1 1 1 1 1 1 1 1 ...
```

The dataset has no missing values:

```
sum(is.na(data))
```

```
## [1] 0
```

For the final dataset used for prediction, the *time* feature (time in days after the patient was dismissed or died) is not selected since this information will not be available for real patients. Binary variables are converted to factors and numeric variables are scaled to improve the prediction. The outcome values' (death event) levels are converted to “No”/“Yes” as needed for some methods.

```
library(tidyverse)
data <- select(data, -time)
data_num <- data
factors <- c(2, 4, 6, 10, 11, 12)
data[factors] <- lapply(data[factors], factor)
data_sc <- data
data_sc[-factors] <- sapply(data[-factors], scale)
levels(data_sc$DEATH_EVENT) = c("No", "Yes")
```

Based on the resulting dataet, train and test sets are created. It is shown that the both sets have a similar proportion of positive and negative outcomes.

```
library(caret)
set.seed(2)
train_index <- createDataPartition(data$DEATH_EVENT, p = 0.8, list = FALSE)
train_not_sc <- train <- data[train_index,]
train <- data_sc[train_index,]
test <- data_sc[-train_index,]
prop.table(table(train$DEATH_EVENT))
```

```
##
##           No           Yes
## 0.6791667 0.3208333
```

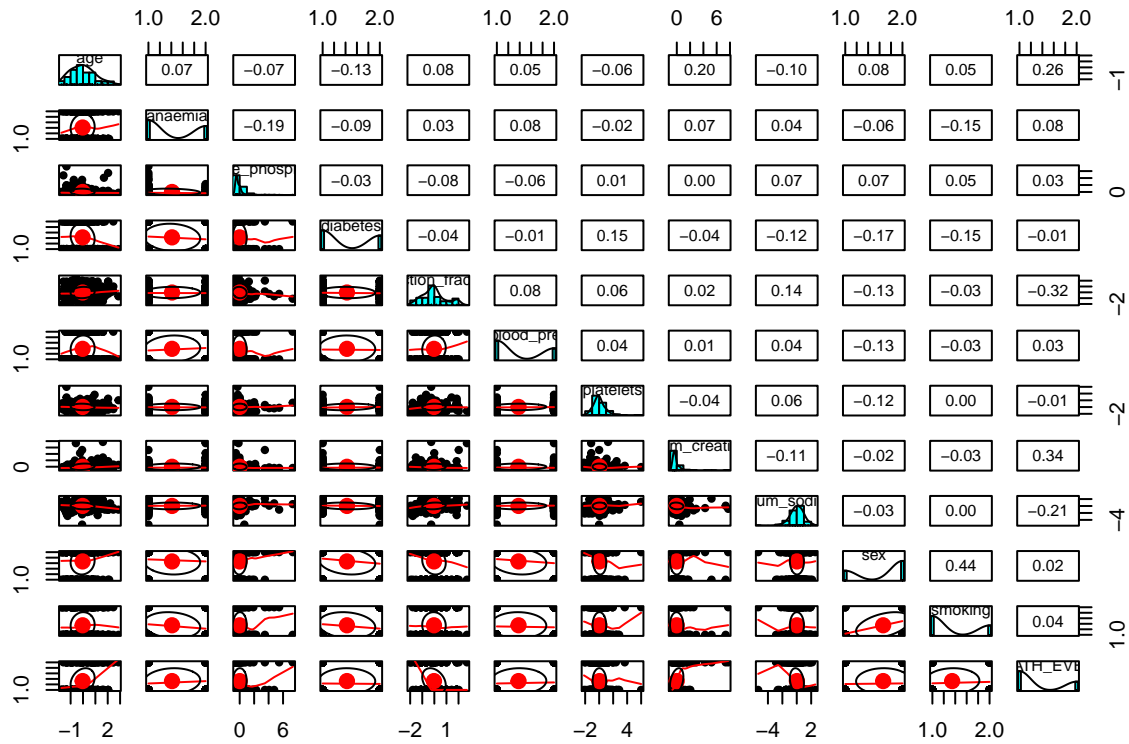
```
prop.table(table(test$DEATH_EVENT))
```

```
##
##           No           Yes
## 0.6779661 0.3220339
```


The train set will be used for the following data exploration.

2.2 Data exploration

The prepared train set contains 6 binary and 5 numeric features and one binary outcome value with 240 observations. The following summary table shows distributions of and correlations between all values and allows a swift examination of the main dependencies.




The following heatmap shows correlations between all values and allows a swift examination of the main dependencies.




As can be seen, the outcome is most strongly correlated with age, ejection fraction, serum creatine and serum sodium features. There is no strong correlation between features except between sex and smoking. All features are evaluated in more detail in the following.

2.2.1 Age

Higher age is expected to be associated with higher mortality. It can be shown in the following histogram with overlaid density curves. The older patients died more often than the younger.

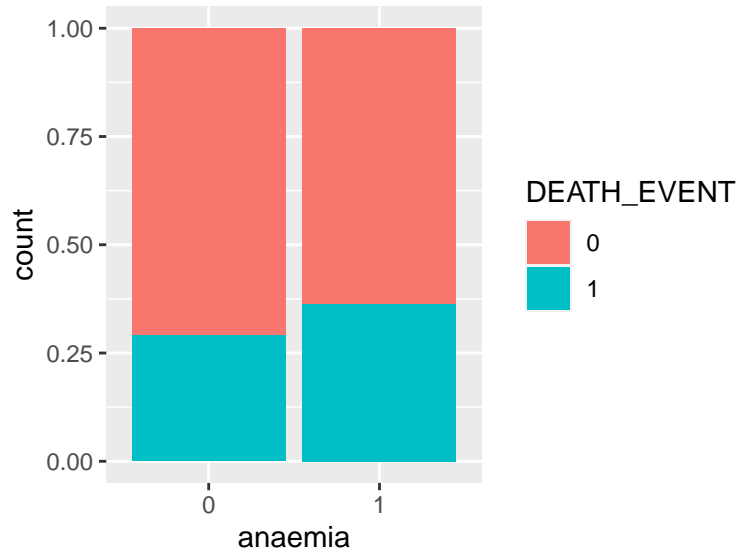


The following boxplots further confirm the assumption.



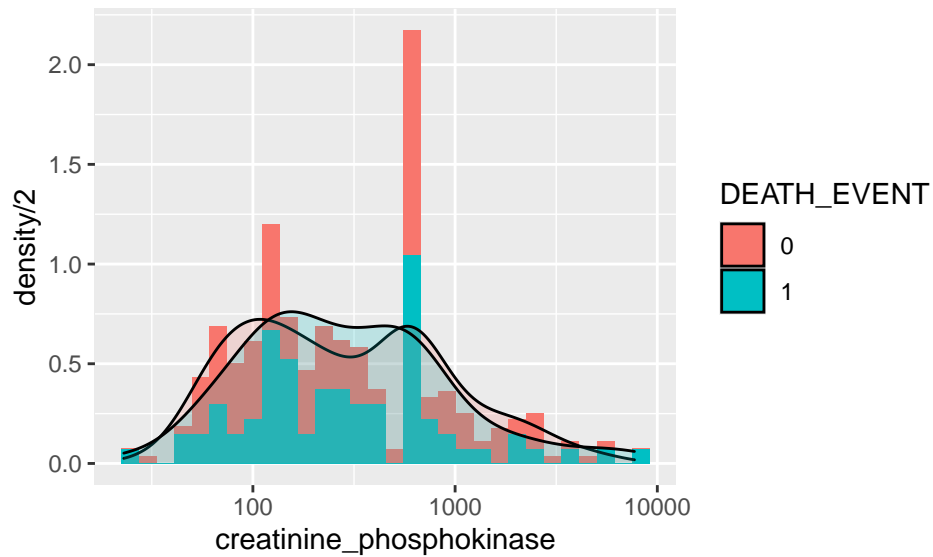
2.2.2 Anaemia

Patients that die seem to have anaemia more often than patients that are dismissed from hospital, as shown in the following percent stacked barchart. The difference is rather small, as expected based on the correlation heatmap.



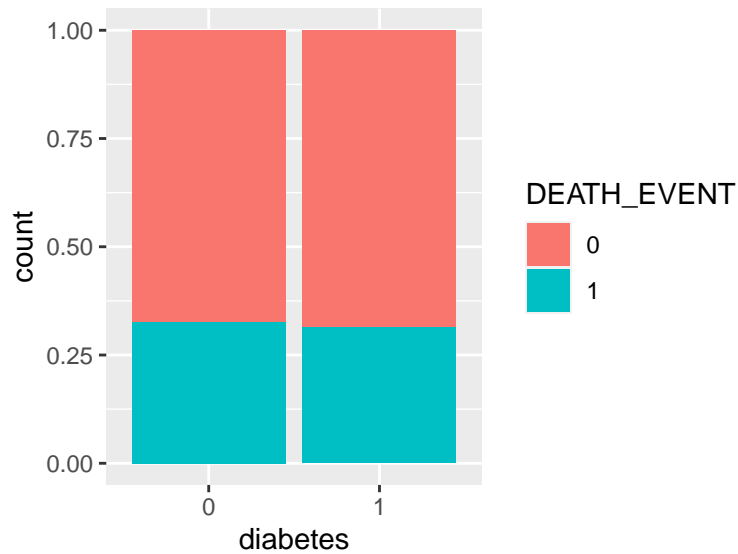
2.2.3 Creatinine phosphokinase

Creatinine phosphokinase levels are similar for both patients groups as shown in the following histogram with overlaid density curves.



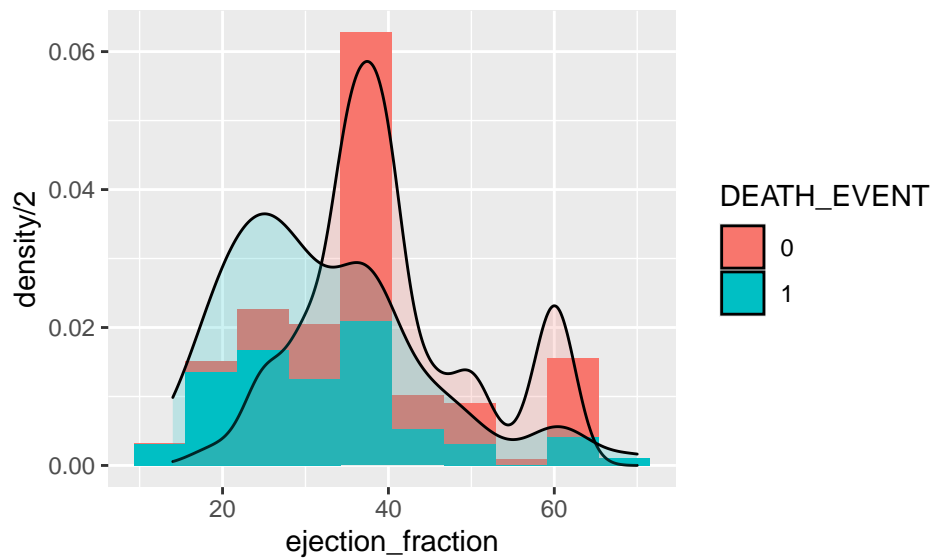
2.2.4 Diabetes

As expected based on the correlation heatmap, there is no evidence that patients suffering from diabetes are more likely to die from heart failure than patients without diabetes for the current dataset. This can be seen in the following percent stacked barchart.



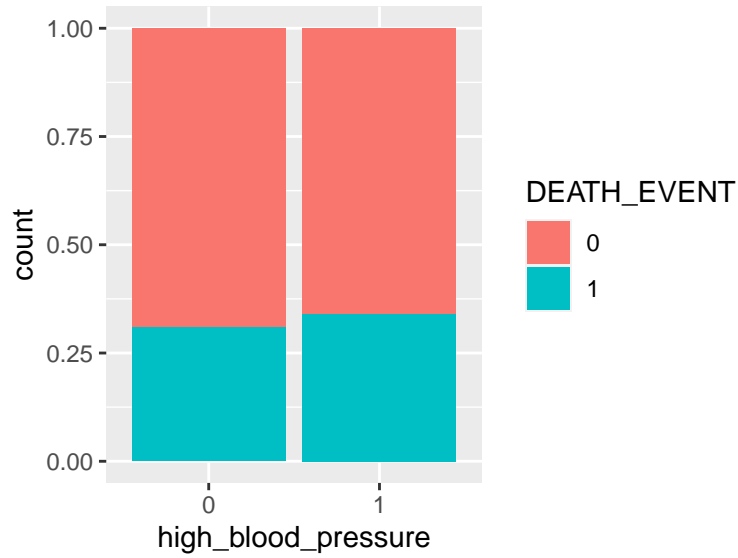
2.2.5 Ejection fraction

Ejection fraction strongly correlates with the outcome. It can be shown in the following histogram with overlaid density curves, where the death events density is skewed to the lower ejection fraction values.



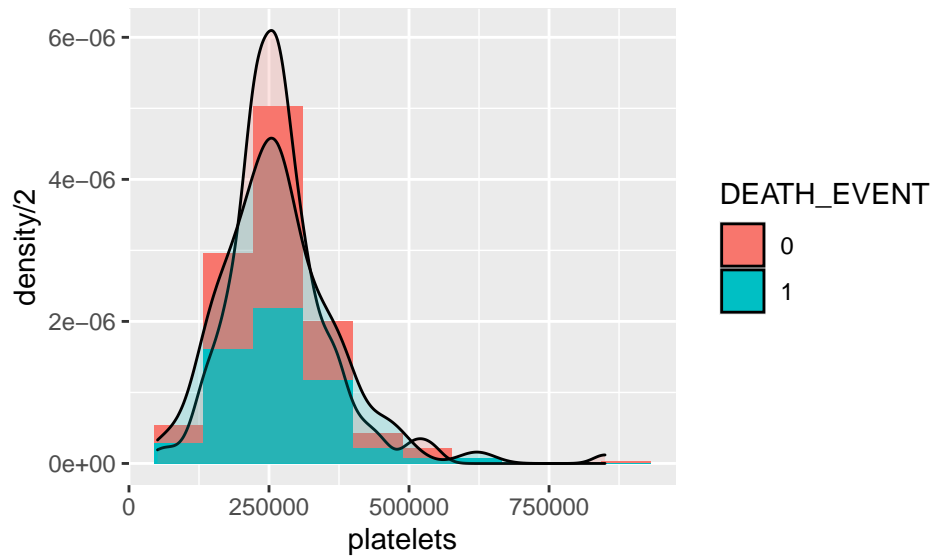
2.2.6 High blood pressure

Patients that die seem to have high blood pressure more often than patients that are dismissed from hospital, as shown in the following percent stacked barchart. The difference is rather small, as expected based on the correlation heatmap.



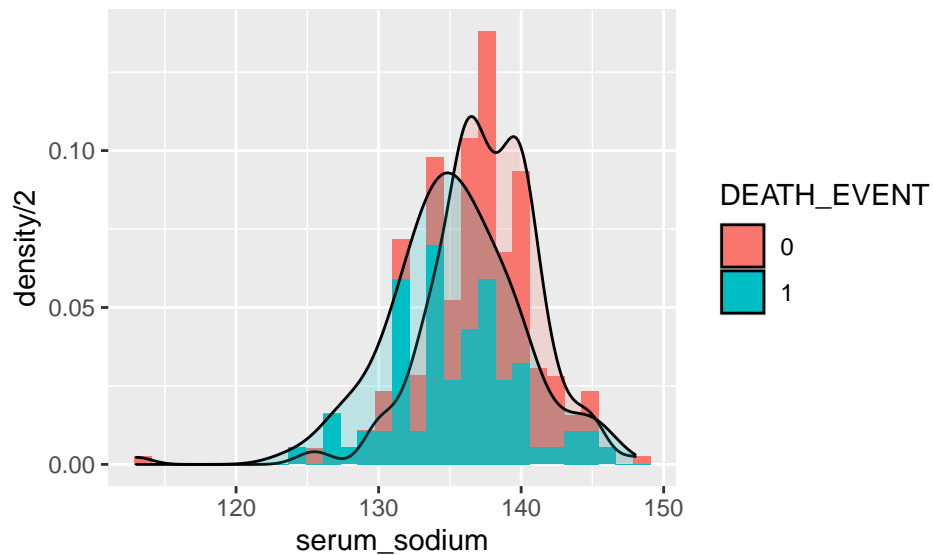
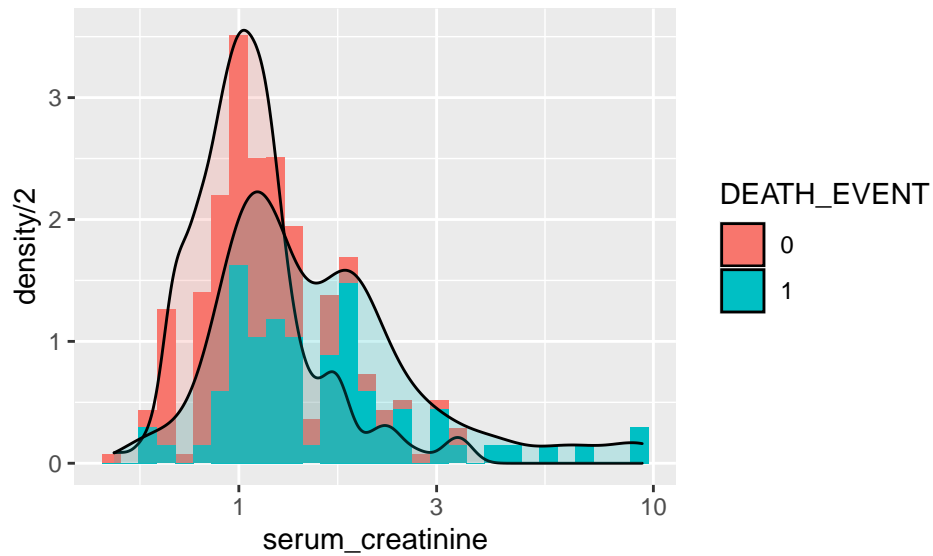
2.2.7 Platelets

There is no difference in platelets levels for both patient groups, as shown in the following histogram with overlaid density curves.



2.2.8 Serum creatinine and Serum sodium

For serum creatinine and serum sodium considerable correlation was shown in the heatmap. The following histograms with overlaid density curves confirm that higher levels of creatinine in sserum and lower levels of sodium in serum correspond to more often cases of patient mortality.




2.2.9 Sex and smoking


Sex and smoking were shown to have a low correlation with patients mortality. As can be seen, there is no significant difference between men and women to die from heart failure, whereas between smoking patients there were more death cases than between non-smoking patients, what can be generally expected.

```
##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine
```



Interestingly, there is a high correlation between sex and smoking, as shown in the following plot. Based on it, one might deduce which sex is presented as “0” and which as “1”.



2.2.10 Findings

2.3 Algorithms

Prediction of patients' mortality (binary outcome) based on several predictors presents a typical classification problem. Machine learning algorithms that will be used in this project are as following:

- Naive Bayes
- Logistic regression
- K-nearest neighbors
- Support vector machine (SVM)
 - Linear
 - Radial
 - Polynomial
- Decision tree
- Bagged decision tree
- Boosted decision tree
- Random forest
- Artificial neural network (ANN)

The algorithms were chosen to represent simple as well as more sophisticated methods. Their performance will be compared based on the selected metric. Additionally, an ensemble model considering prediction of all said models will be tried out.

2.4 Performance evaluation

As it was shown before, the data set includes about 70% of negative outcomes and 30% of positive outcomes. This means that the data is imbalanced, and it is not recommended to use the accuracy metric for model evaluation, since it might show relatively high values for obviously bad predictions. For example, should all outcomes in the train set be predicted as negative, the accuracy would still be as high as 0.683:

```
pred_neg <- c(1, rep(0, nrow(train) - 1))
pred_neg <- factor(pred_neg)
levels(pred_neg) = c("No", "Yes")
cm_neg <- confusionMatrix(pred_neg, train$DEATH_EVENT, mode = "everything", positive = "Yes")
cm_neg$overall["Accuracy"]
```

```
## Accuracy
## 0.6833333
```

Alternatively, F1-score will be used as suggested by numerous sources (e.g. as summarized in [this article] (<https://towardsdatascience.com/metrics-for-imbalanced-classification-41c71549bbb5>)). Death event will be used as a “positive” outcome since it is more important for the prediction — the cost of the mistake might be patients death. A false positive is less crucial in this sense. For the above example, the F1-score is only 0.0256:

```
cm_neg$byClass["F1"]
```

```
## F1
## 0.02564103
```

10-fold cross-validation will be used for model training. The final model evaluation will be performed with the test set.

3 Results

3.1 Preparations

caret package will be used for model training and evaluation. Since there is no build-in function for training on F1-score, the necessary function is written first.

```
f1 <- function(data, lev = NULL, model = NULL) {  
  f1_val <- MLmetrics::F1_Score(y_pred = data$pred,  
                                y_true = data$obs,  
                                # level 2 ("Yes") will be used as positive level  
                                positive = lev[2])  
  
  c(F1 = f1_val)  
}
```

With the *trainControl* function the training method — 10-fold cross-validation — and the metric — F1-score — are defined:

```
fitControl <- trainControl(method = "cv",  
                           number = 5,  
                           p = 0.8,  
                           classProbs = TRUE,  
                           summaryFunction = f1)
```

correct!!! The following function will be used to quickly try several algorithms. It will take a vector of methods to train and a list algorithm parameters and will return a list containing trained models, confusion matrices and a dataframe with selected metrics calculated for the train set (F1-score, accuracy and specificity).

```
calc_model <- function(method, tuneGrid) {  
  set.seed(3)  
  fit <- train(DEATH_EVENT ~., # all features will be considered first  
              data = train,  
              method = method,  
              metric = "F1",  
              trControl = fitControl,  
              tuneGrid = tuneGrid)  
  
  # F1-score for the train set will be used to evaluate models prior to evaluation  
  # based on the test set  
  F1_train <- max(fit$results$F1, na.rm = TRUE)  
  # Prediction for the train set  
  pred <- predict(fit, train)  
  cm <- confusionMatrix(pred, train$DEATH_EVENT, mode = "everything", positive = "Yes")  
  accuracy <- cm$overall["Accuracy"]  
  F1 <- cm$byClass["F1"]  
  results <- list()  
  # The functions returns a list of trained models, confusion matrices and a dataframe with metrics  
  results[[1]] <- fit  
  results[[2]] <- cm  
  results[[3]] <- data.frame(method = method, F1_train = F1_train, F1 = F1, accuracy = accuracy)  
  return(results)  
}
```

The following list includes all tuning parameters for training. Later in the chapter graphical representation of tuning will be used to reassess the choice of parameters.

```
tuneGrids <- list()

tuneGrids[[1]] <- expand.grid(fL = seq(0, 5, 1),
                             usekernel = c(TRUE, FALSE),
                             adjust = seq(0, 5, 1))

tuneGrids[2] <- list(NULL)

tuneGrids[[3]] <- data.frame(k = seq(1, 99, 2))

tuneGrids[[4]] <- data.frame(C = seq(0.5, 10, 0.5))

tuneGrids[5] <- list(NULL)

tuneGrids[6] <- list(NULL)

tuneGrids[[7]] <- data.frame(cp = seq(0, 0.05, len = 10))

tuneGrids[8] <- list(NULL)

tuneGrids[[9]] <- data.frame(nIter = seq(1, 19, 2),
                             method = "M1")

tuneGrids[[10]] <- data.frame(mtry = seq(1, 11, 1))

tuneGrids[[11]] <- expand.grid(size = seq(1, 10, 1),
                              decay = seq(1, 10, 1))
```

3.2 Model training

The models are trained using the defined parameters and the resulting metrics are saved as a dataframe. The F1-score obtained in the cross-validation is opposed to the F1-score and accuracy calculated using the trained model for the train set.

```
methods <- c("nb", "glm", "knn", "svmLinear", "svmRadial", "svmPoly", "rpart",
             "treebag", "adaboost", "rf", "nnet")

# wrapping the function in capture.output prevents printing of train() console messages
silent <- capture.output(
  all_models <- mapply(calc_model, methods, tuneGrids))

df <- bind_rows(all_models[seq(3, 33, 3)], .id = "column_label")
rownames(df) <- NULL
df <- df[, -1]
knitr::kable(df)
```

method	F1_train	F1	accuracy
nb	0.4652814	0.6218487	0.8125000

method	F1_train	F1	accuracy
glm	0.6056817	0.6666667	0.8083333
knn	0.4918758	0.5737705	0.7833333
svmLinear	0.5584230	0.6412214	0.8041667
svmRadial	0.6075115	0.7248322	0.8291667
svmPoly	0.6385619	0.6575342	0.7916667
rpart	0.5659306	0.7297297	0.8333333
treebag	0.5211879	1.0000000	1.0000000
adaboost	0.5123990	0.9241379	0.9541667
rf	0.5837108	1.0000000	1.0000000
nnet	0.5664762	0.6212121	0.7916667

As can be seen, the F1-score values of the cross-validation are lower for all methods, especially for the bagged and boosted classification trees and random forest algorithms, what indicates overtraining. In order to overcome this, we will try reduce the number of predictors leaving for training only features that showed correlation with the outcome value more than 0.05 (refer to section 2.2). The function is amended and models are trained anew.

```
calc_model_red <- function(method, tuneGrid) {
  set.seed(3)
  fit <- train(DEATH_EVENT ~ age + anaemia + ejection_fraction + serum_creatinine +
    serum_sodium, # only features with correlation >= 0.05 are left
    data = train,
    method = method,
    metric = "F1",
    trControl = fitControl,
    tuneGrid = tuneGrid)
  F1_train <- max(fit$results$F1, na.rm = TRUE)
  pred <- predict(fit, train)
  cm <- confusionMatrix(pred, train$DEATH_EVENT, mode = "everything", positive = "Yes")
  accuracy <- cm$overall["Accuracy"]
  F1 <- cm$byClass["F1"]
  results <- list()
  results[[1]] <- fit
  results[[2]] <- cm
  results[[3]] <- data.frame(method = method, F1_train = F1_train, F1 = F1, accuracy = accuracy)
  return(results)
}

silent <- capture.output(
  all_models_red <- mapply(calc_model_red, methods, tuneGrids))

df_red <- bind_rows(all_models_red[seq(3, 33, 3)], .id = "column_label")
rownames(df_red) <- NULL
df_red <- df_red[, -1]
knitr::kable(df_red)
```

method	F1_train	F1	accuracy
nb	0.5657867	0.6615385	0.8166667
glm	0.6156242	0.6376812	0.7916667
knn	0.5656920	0.5909091	0.7750000

method	F1_train	F1	accuracy
svmLinear	0.5598537	0.6165414	0.7875000
svmRadial	0.6569020	0.7417219	0.8375000
svmPoly	0.6488063	0.6447368	0.7750000
rpart	0.6084977	0.7341772	0.8250000
treebag	0.5561290	0.9935484	0.9958333
adaboost	0.5809217	1.0000000	1.0000000
rf	0.6039143	1.0000000	1.0000000
nnet	0.5637738	0.6259542	0.7958333

A certain improvement of F1-score values can be observed. Additionally, the difference between the F1-score values calculated within the cross-validation and using the trained model for the train set became smaller what means less overtraining. Therefore, the reduced predictors set will be kept for the further analysis.

The separate models will be evaluated in the following sections.


3.2.3 Naive Bayes

The simple algorithm Naive Bayes shows surprisingly good performance in comparison with the other models. The tuning plot shows that no further prediction improvement through choosing parameters is possible.

```
all_models_red[[2]]$table
```

```
##           Reference
## Prediction  No Yes
##           No 153 34
##           Yes  10 43
```

```
plot(all_models_red[[1]])
```



```
all_models_red[[3]]$F1_train
```

```
## [1] 0.5657867
```

3.2.4 Logistic regression

Logistic regression showed a slightly better performance than Naive Bayes. There are no model parameters to tune.

```
all_models_red[[6]]$F1_train
```

```
## [1] 0.6156242
```

```
all_models_red[[5]]$table
```

```
##           Reference
## Prediction No Yes
##      No  146  33
##      Yes   17  44
```

3.2.5 K-nearest neighbors

K-nearest neighbors show similar performance to Naive Bayes. Higher number of neighbors seem only to decrease the F1-score.


```
all_models_red[[9]]$F1_train
```

```
## [1] 0.565692
```

```
all_models_red[[8]]$table
```

```
##           Reference
## Prediction No Yes
##      No  147  38
##      Yes   16  39
```

```
plot(all_models_red[[7]])
```

3.2.6 SVM Linear

SVM Linear shows performance similar to Logistic regression. Higher cost values can be tried out in order to increase the F1-score.

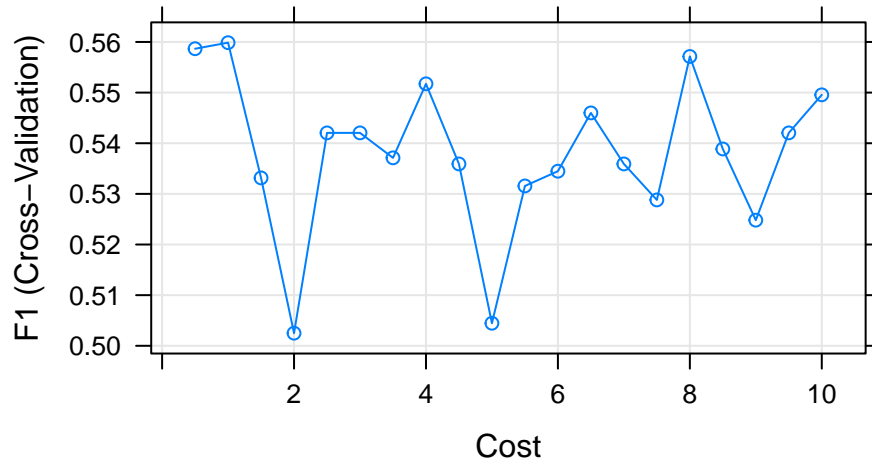
```
all_models_red[[12]]$F1_train
```

```
## [1] 0.5598537
```

```
all_models_red[[11]]$table
```

```
##           Reference
## Prediction No Yes
##           No 148 36
##           Yes  15 41
```

```
plot(all_models_red[[10]])
```



3.2.7 SVM Radial

SVM Radial shows the highest performance of all trained models. Further improvement through parameters tuning is, however, not possible.


```
all_models_red[[15]]$F1_train
```

```
## [1] 0.656902
```

```
all_models_red[[14]]$table
```

```
##           Reference
## Prediction No Yes
##           No 145 21
##           Yes  18 56
```

```
plot(all_models_red[[13]])
```



3.2.8 SVM Polynomial

SVM Polynomial shows performance close to the SVM Radial model. Also here, the tuning plot shows that no further prediction improvement through choosing parameters is possible.


```
all_models_red[[18]]$F1_train
```

```
## [1] 0.6488063
```

```
all_models_red[[17]]$table
```

```
##           Reference
## Prediction No Yes
##           No 137 28
##           Yes 26 49
```

```
plot(all_models_red[[16]])
```



3.2.9 Decision tree

Decision tree algorithm shows performance close to the SVM Radial model. No further improvement through parameters tuning is possible.


```
all_models_red[[21]]$F1_train
```

```
## [1] 0.6084977
```

```
all_models_red[[20]]$table
```

```
##           Reference
## Prediction No Yes
##          No 140 19
##          Yes 23 58
```

```
plot(all_models_red[[19]])
```



3.2.10 Bagged decision tree

Bagged decision tree shows surprisingly low performance similar to Naive Bayes. The confusion table indicates strong overtraining. There are no model parameters to tune.

```
all_models_red[[24]]$F1_train
```

```
## [1] 0.556129
```

```
all_models_red[[23]]$table
```

```
##           Reference
## Prediction  No Yes
##           No 162  0
##           Yes  1  77
```

3.2.11 Boosted decision tree

Relatively low performance in cross-validation (similar Naive Bayes) opposed to rather good classification of the train set indicates overtraining in this case. Higher number of trees will be tried out in order to improve performance and minimize overtraining.


```
all_models_red[[27]]$F1_train
```

```
## [1] 0.5809217
```

```
all_models_red[[26]]$table
```

```
##           Reference
## Prediction  No Yes
##           No 163  0
##           Yes  0  77
```

```
plot(all_models_red[[25]])
```



3.2.12 Random forest

Random forest shows performance similar to the SVM Linear model. The confusion table indicates strong overtraining. A higher number of randomly selected predictors can be used to improve performance/reduce overtraining.


```
all_models_red[[30]]$F1_train
```

```
## [1] 0.6039143
```

```
all_models_red[[29]]$table
```

```
##           Reference
## Prediction No Yes
##      No  163  0
##      Yes   0  77
```

```
plot(all_models_red[[28]])
```



3.2.13 ANN

ANN shows performance similar to the Random forest model. No further improvement through parameters tuning is possible.


```
all_models_red[[33]]$F1_train
```

```
## [1] 0.5637738
```

```
all_models_red[[32]]$table
```

```
##           Reference
## Prediction No Yes
##           No 150 36
##           Yes 13 41
```

```
plot(all_models_red[[31]])
```



3.2.14 Further parameter tuning

The above observations are used to redefine tuning parameters ranges for some models. The influence of the amended parameter ranges is described in the following.

Tuning parameters are redefined for SVM linear, Boosted decision tree and Random forest:

```
tuneGrids[[4]] <- data.frame(C = seq(0.5, 100, 0.5))

tuneGrids[[9]] <- data.frame(nIter = seq(1, 39, 2),
                             method = "M1")

tuneGrids[[10]] <- data.frame(mtry = seq(1, 20, 1))
```

And the models are trained again:

```
silent <- capture.output(
  all_models_red <- mapply(calc_model_red, methods, tuneGrids))
```

The results are checked again.

SVM Linear:

Higher cost values indeed resulted in a slight increase of the F1-score.

```
all_models_red[[12]]$F1_train
```


```
## [1] 0.5694028
```

```
all_models_red[[11]]$table
```

```
##           Reference
## Prediction No Yes
##           No 149 38
##           Yes  14 39
```



```
plot(all_models_red[[10]])
```



Boosted decision tree:

Higher number of trees did not result in performance improvement.

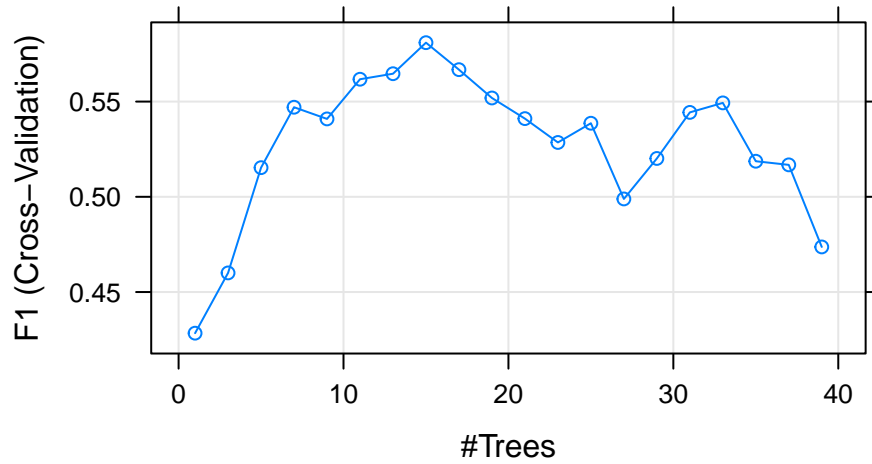
```
all_models_red[[27]]$F1_train
```

```
## [1] 0.5809217
```

```
all_models_red[[26]]$table
```

```
##           Reference
## Prediction  No Yes
##          No 163  0
##          Yes  0  77
```

```
plot(all_models_red[[25]])
```



Random forest:

A higher number of randomly selected predictors did not result in performance improvement.


```
all_models_red[[30]]$F1_train
```

```
## [1] 0.6219926
```

```
all_models_red[[29]]$table
```

```
##           Reference
## Prediction No Yes
##          No 163  0
##          Yes  0  77
```

```
plot(all_models_red[[28]])
```



Thus, by changing the parameters' range it was only possible to improve model performance of the SVM Linear model. The final table with the trained models is presented below.

```
df_red <- bind_rows(all_models_red[seq(3, 33, 3)], .id = "column_label")
rownames(df_red) <- NULL
df_red <- df_red[, -1]
knitr::kable(df_red)
```

method	F1_train	F1	accuracy
nb	0.5657867	0.6615385	0.8166667
glm	0.6156242	0.6376812	0.7916667
knn	0.5656920	0.5909091	0.7750000
svmLinear	0.5694028	0.6000000	0.7833333
svmRadial	0.6569020	0.7417219	0.8375000
svmPoly	0.6488063	0.6447368	0.7750000
rpart	0.6084977	0.7341772	0.8250000
treebag	0.5561290	0.9935484	0.9958333
adaboost	0.5809217	1.0000000	1.0000000
rf	0.6219926	1.0000000	1.0000000
nnet	0.5637738	0.6259542	0.7958333

The best trained model uses SVM Radial algorithm, whereas the other SVM models (Linear and Polynomial) also show relative high performance values. The models K-Nearest neighbor, Bagged decision tree, Boosted decision tree and Random forest showed strong overtraining for the train set, what might be due to the small set size.

3.2.15 Final model evaluation

The test set is now used to make predictions with the best model and to perform the final evaluation. F1-score value is slightly lower for the test set, but still close enough to indicate that no strong overtraining is observed. The confusion matrix shows that the model was possible to predict most of the negative and positive outcomes correctly. As said above, false negative prediction rate is more important than false positive rate in the current case. The models predicts death event only for 12 patients out of 19 actual positive outcomes (63%).

```
svmRadial_pred <- predict(all_models_red[[13]], test)
svmRadial_cm <- confusionMatrix(svmRadial_pred, test$DEATH_EVENT, mode = "everything", positive = "Yes")
svmRadial_cm$byClass["F1"]
```

```
##          F1
## 0.6190476
```

```
svmRadial_cm$table
```

```
##          Reference
## Prediction No Yes
##          No  30   6
##          Yes 10  13
```

3.3 Ensemble learning

Ensemble learning can be used to improve predictions by considering results of all trained models. The predictions of all models are evaluated like votes, and either the majority of votes or number of votes above a threshold are considered as a positive outcome. To apply this approach, first, we will use the following function that returns a dataframe with predicted outcomes for the train set.

```
pred_all <- function(method, tuneGrid) {
  set.seed(3)
  fit <- train(DEATH_EVENT ~ age + anaemia + ejection_fraction + serum_creatinine +
    serum_sodium,
    data = train,
    method = method,
    metric = "F1",
    trControl = fitControl,
    tuneGrid = tuneGrid)
  pred <- predict(fit, train)
  df <- data.frame(method = pred)
  return(df)
}
```

The function creates a dataframe with “Yes”/“No” populated column for each model. The models K-Nearest neighbors, Bagged decision tree, Boosted decision tree and Random forest are excluded from the ensemble approach since they showed strong overtraining, as shown before. Factor output “Yes”/“No” is confirmed to numeric 1/0 and the row sums are used as a cumulative prediction.

```
silent <- capture.output(
  all <- mapply(pred_all, methods, tuneGrids))
all <- as.data.frame(all)
all <- all[, -c(3, 8, 9, 10)] # Models with strong overtraining are excluded

# Output values are transformed to 1s and 0s
all <- ifelse(all == "Yes", 1, 0)
# Summary prediction is calculated
all_sums <- rowSums(all)
all_sums
```

```
## [1] 7 7 7 7 0 7 7 0 1 0 1 2 6 5 6 7 0 0 7 3 0 7 7 7 1 6 3 6 3 7 1 0 0 2 1 0 7
## [38] 0 6 7 0 6 7 0 2 7 1 2 0 0 7 7 5 7 0 0 5 0 7 2 0 0 0 0 1 0 7 3 1 0 0 0 0 0
## [75] 0 0 0 7 0 0 0 5 0 2 0 4 0 1 7 0 0 0 2 0 4 0 1 0 0 3 0 7 0 0 0 0 7 0 5 0 0
## [112] 0 3 0 7 0 0 2 0 0 0 7 0 0 0 4 0 0 0 0 5 4 0 0 0 0 1 3 0 7 0 0 0 0 0 0 0 0
## [149] 0 6 4 5 0 7 0 7 0 0 3 0 6 4 0 0 0 6 0 0 2 0 0 7 0 4 0 0 7 1 0 7 0 0 1 0 0
## [186] 7 0 0 0 5 0 2 0 0 1 0 1 7 0 0 0 0 0 1 0 0 0 0 0 0 0 0 7 0 0 7 0 0 4 0 0 0
## [223] 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 7 0 0 7 0 0 4 0 0 0
```

In order to find out, which threshold should be used to separate positive and negative outcomes, the function `ensemble()` is used. It takes a vote threshold needed for a positive or negative vote and returns the corresponding F1-score value. Vote thresholds 1–7 are tried and corresponding F1-scores are calculated. Finally, the relationship between the threshold and F1-score is shown in the plot.

```
ensemble <- function(vote) {
  all_voted <- ifelse(all_sums >= vote, 1, 0)
  all_voted <- as.factor(all_voted)
```


```

levels(all_voted) = c("No", "Yes")
cm_voted <- confusionMatrix(all_voted, train$DEATH_EVENT, positive = "Yes")
return(cm_voted$byClass["F1"])
}

vote <- seq(1, 7, 1)
F1_ens <- sapply(vote, ensemble)

df_ens <- data.frame(vote = vote, F1_ens = F1_ens)
df_ens %>%
  ggplot(aes(vote, F1_ens)) +
    geom_line() +
    xlab("Vote threshold") +
    ylab("F1 ensemble")

```



The best threshold is just 1, what means that if any of the considered models predicts a positive outcome, then the ensemble prediction is also a positive outcome. In other words, it is like a full outer join of all predictions. The F1-score is slightly higher than this of all included models (calculated for the train set, not in cross-validation).

```

all_voted <- ifelse(all_sums >= 1, 1, 0)
all_voted <- as.factor(all_voted)
levels(all_voted) = c("No", "Yes")
cm_voted <- confusionMatrix(all_voted, train$DEATH_EVENT, positive = "Yes")
cm_voted$byClass["F1"]

```

```

##          F1
## 0.7471264

```

Though, the use of threshold = 1 doesn't look to be very helpful, the performance of the ensemble prediction is still evaluated.

```

pred_all_test <- function(method, tuneGrid) {
  set.seed(3)
  fit <- train(DEATH_EVENT ~ age + anaemia + ejection_fraction + serum_creatinine +
    serum_sodium,
    data = train,
    method = method,
    metric = "F1",
    trControl = fitControl,
    tuneGrid = tuneGrid)
  pred <- predict(fit, test)
  df <- data.frame(method = pred)
  return(df)
}

silent <- capture.output(
  all_test <- mapply(pred_all_test, methods, tuneGrids))
all_test <- as.data.frame(all_test)
all_test <- all_test[,-c(3, 8, 9, 10)]

all_test <- ifelse(all_test == "Yes", 1, 0)
all_test_sums <- rowSums(all_test)

all_test_voted <- ifelse(all_test_sums >= 1, 1, 0)
all_test_voted <- as.factor(all_test_voted)
levels(all_test_voted) = c("No", "Yes")
cm_test_voted <- confusionMatrix(all_test_voted, test$DEATH_EVENT, positive = "Yes")
cm_test_voted$byClass["F1"]

##          F1
## 0.5714286

```

As can be seen, the performance of the ensemble is lower than this of the best trained model.

4 Conclusion

In this project, we have trained and evaluated several machine learning models to predict mortality caused by heart failure based on patients features.

The data set was analysed and evaluated in order to find important data relationships and to select meaningful predictors for the model training. The data was prepared for the training, whereas numeric predictors were scaled. F1-score was chosen to evaluate model performance based on the imbalanced nature of data. Different algorithms were used then to train the models and the best algorithm was identified.

Radial Support Vector Machine (SVM Radial) showed the highest performance of all models in the 10-fold cross-validation (F1-score: 0.66). The models predicts death event only for 12 patients out of 19 actual positive outcomes (63%), what would be insufficient in the real world. The performance might be improved by using a larger training data set, 299 observations is modest.

The other tried algorithms showed lower performance and some of them showed a strong overtraining. The attempts to improve the performance/to reduce overtraining were fruitless though. The application of the ensemble learning approach also didn't result in any improvement.

4.1 Limitations

As already mentioned, the small size of the available data set was the main limitation in the project.

4.2 Outlook

The selection of algorithms used to train the models was chosen to include simple as well as more sophisticated methods of different types. For the future work, a more accurate selection based on an extensive literature research can be performed. Algorithms better tailored for small data sets might help to improve the prediction.